

Soon after the FLP impossibility result appeared, people started to find a way to circumvent it. In 1983, Ben Or gave the first *randomized algorithm* that solve the consensus problem in asynchronous message-passing system with $f < n/2$ crashes with probability 1 (n is the number of processes and f is the maximum number of processes that can crash). The protocol runs forever, but every non-faulty process makes a decision in finitely many steps. With high probability this number of steps is huge ($O(2^{2^n})$). In spite of the fact that Ben Or's protocol is not practical the result is important because it indicates that by using randomization the impossibility of consensus can be circumvented.

The model of the distributed system is almost the same: we have a set of n processes, exchange messages in an asynchronous environment. Communication channels are reliable. Each process has access to a *perfect coin*: when a process toss its coin, it obtains 0 or 1 with probability $1/2$. The consensus protocol must satisfy the usual conditions, except for the Termination Property, that is slightly relaxed:

- **Integrity:** At most one decision per process.
- **Agreement:** All correct processes must choose the same decision value.
- **Validity:** The value decided by a process must have been initially proposed.
- **(New) Termination:** *With probability 1*, every correct process eventually decides some value.

Every process p_i executes the **randomized consensus algorithm**, as reported in Algorithm 1.

The algorithm works under the assumption that a majority of processes are correct (i.e., $n > 2f$). It is easy to see that this requirement is necessary for any algorithm that solves the consensus problem in asynchronous systems with crash failures, even if all processes have access to a random number generator.

In the algorithm, every message contains a tag (R or P), a round number k , and a value which is either 0 or 1 (for messages tagged P, it could also be ?). Messages tagged R are called *reports* and those tagged with P are called *proposals*. When p_i sends (R, k , v) or (P, k , v) we say that p_i reports or proposes v in round k , respectively.

Each execution of the **while** loop is called a *round*, and each round consists of two asynchronous *phases*. During round k only messages which are timestamped with round number k are processed.

In the first phase, processes report to each other their current estimate (0 or 1) for a decision value.

In the second phase, if a process receives a majority of reports for the same value then it proposes that value to all processes, otherwise it proposes "?". Note that it is impossible for one process to propose 0 and another process to propose 1 in the same round. At the end of the second phase, if a process receives $f + 1$ proposals for the same value different than ?, then it decides that

Algorithm 1: The *randomized consensus* algorithm

Input: The initial value v_i of process p_i

RANDOMIZEDCONSENSUS(v_i):

```
1   $x \leftarrow v_i$  //  $p_i$ 's current estimate of the decision value
2   $k \leftarrow 0$  //  $k$  is the current round number
3  while true do
4     $k \leftarrow k + 1$ 
5    SEND ( $R, k, x$ ) to all processes
6    WAIT for ( $R, k, *$ ) messages from  $n - f$  processes ( $* \in \{0, 1\}$ )
7    if received more than  $n/2$  ( $R, k, v$ ) messages with the same  $v$  then
8      SEND ( $P, k, v$ ) to all processes
9    else
10     SEND ( $P, k, ?$ ) to all processes
11    WAIT for ( $P, k, *$ ) messages from  $n - f$  processes ( $* \in \{0, 1, ?\}$ )
12    if received at least  $f + 1$  ( $P, k, v$ ) messages with the same  $v \neq ?$  then
13      decide  $v$ 
14    if received at least one ( $P, k, v$ ) message with  $v \neq ?$  then
15       $x \leftarrow v$ 
16    else
17       $x \leftarrow 0$  or  $1$  randomly // toss coin
```

value. If it receives at least one value different than $?$, then it adopts that value as its new estimate, otherwise it adopts a random value for its estimate.

The algorithm does not include a halt statement. Moreover, once a correct process decides a value, it will keep deciding the same value in all subsequent phases. However, it is easy to modify the algorithm so that every process decides at most once, and halts at most one phase after deciding.

The Integrity Property is guaranteed by the semantics of the **decide** primitive at line 12. We must show that this algorithm satisfies the remaining three conditions.

We now give some definitions.

- We say that *process p_i starts round k* if process p_i completes at least $k - 1$ iterations of the **while** loop.
- We say that *process p_i reaches line n in round k* if process p_i starts round k and p_i executes past line $n - 1$ in that round.
- We say that *value v is k -locked* if every process that starts round k does so with its variable x set to v .

When ambiguities may arise, a local variable of a process p_i is subscripted by p_i , e.g., x_{p_i} is the local variable x of process p_i .

Lemma 3. *If it is impossible for process p_i to propose 0 and for process p_j to propose 1 in the same round $k > 0$.*

Proof. We proceed by contradiction. Suppose that process p_i proposes 0 and process p_j proposes 1 in round k . Thus p_i receives more than $n/2$ reports equal to 0 and p_j receives more than $n/2$ reports equal to 1 in round k . Thus it exists a process p_k that reports 0 to p_i and 1 to p_j in round k , but this is impossible. \square

Lemma 4. *If some process p_i decides v in round $k > 0$, then all the processes p_j that start round $k + 1$ do so with $x_{p_j} = v$, i.e., v is $k + 1$ -locked.*

Proof. Suppose that some process p_i decides v in round k . Process p_i must have received at least $f + 1$ proposals for v in round k . Let p_j be any process that starts round $k + 1$: then p_j received $n - f$ proposals in round k . Hence p_j receives at least one v in round k since $n - f > f + 1$. By the previous Lemma 3, p_j did not receive \bar{v} in round k , hence p_j sets its $x = v$ in round k and p_j starts round $k + 1$ with $x_{p_j} = v$. \square

Lemma 5. *If a value v is k -locked for some $k > 0$, then every process that reaches line 13 in round k decides v in round k .*

Proof. Suppose v is k -locked for some $k > 0$. Then, all reports sent in line 6 of round k are for v . Since $n - f > n/2$, every process that proposes some value in round k proposes v in line 8. Consider a process p that reaches line 13 in round k . Clearly, p receives $n - f$ proposals for v in round k . Since $n - f \geq f + 1$, p decides v in round k . \square

Corollary 1. *If some process decides v in round $k > 0$, then every processes that executes line 13 in round $k + 1$ decides v in round $k + 1$.*

Proof. From Lemma 4 and Lemma 5. \square

Corollary 2 (Agreement). *If some processes p_i and p_j decide v and v' in round $k > 0$ and $k' > 0$, respectively, then $v = v'$.*

Proof. For $k = k'$ the result follows from Lemma 3 and the fact that a process can decide a value in a round only if that value was proposed in the same round. Assume that $k < k'$. Since p' decides in round k' then p' reaches line 13 in every round r , $k < r \leq k'$. Since p decides v in round k , by Corollary 1 p' decides v in round $k + 1 \leq k'$. By additional applications of Corollary 1, we conclude that p' decides v in round k' . Each process can decide at most once per round, so $v = v'$. \square

Corollary 3 (Validity). *If some process p decides v , then v is the initial value of some process.*

Proof. Note that $v \in \{0, 1\}$. If the initial values of all processes are not identical, then v is clearly the initial value of some process. Now, suppose all processes have the same initial value w . Thus, w is 1-locked. From Lemma 5, p decides w , and from Corollary 2, $w = v$. \square

Theorem ((New) Termination).

$$\lim_{k \rightarrow \infty} \text{P}[v \text{ is } k\text{-locked during the first } k \text{ rounds}] = 1.$$

Proof. By Lemma 5, if some value v is k -locked, then v is decided in round k . At round 1,

$$\text{P}[v \text{ is 1-locked during the first round}] = \frac{1}{2^n}.$$

At round k , the probability that some value v is k -locked is *at least* $\frac{1}{2^n}$; indeed some process p_i can set $x_i = v$ not necessarily by flipping a coin. Hence, for any round k ,

$$\text{P}[\text{no value is } k\text{-locked}] < 1 - \frac{1}{2^n}.$$

Since coin flips are independent,

$$\text{P}[\text{no value is } k\text{-locked for the first } k \text{ rounds}] < \left(1 - \frac{1}{2^n}\right)^k.$$

Thus the probability that v is k -locked during the first k rounds is

$$\text{P}[v \text{ is } k\text{-locked for the first } k \text{ rounds}] \geq 1 - \left(1 - \frac{1}{2^n}\right)^k.$$

□