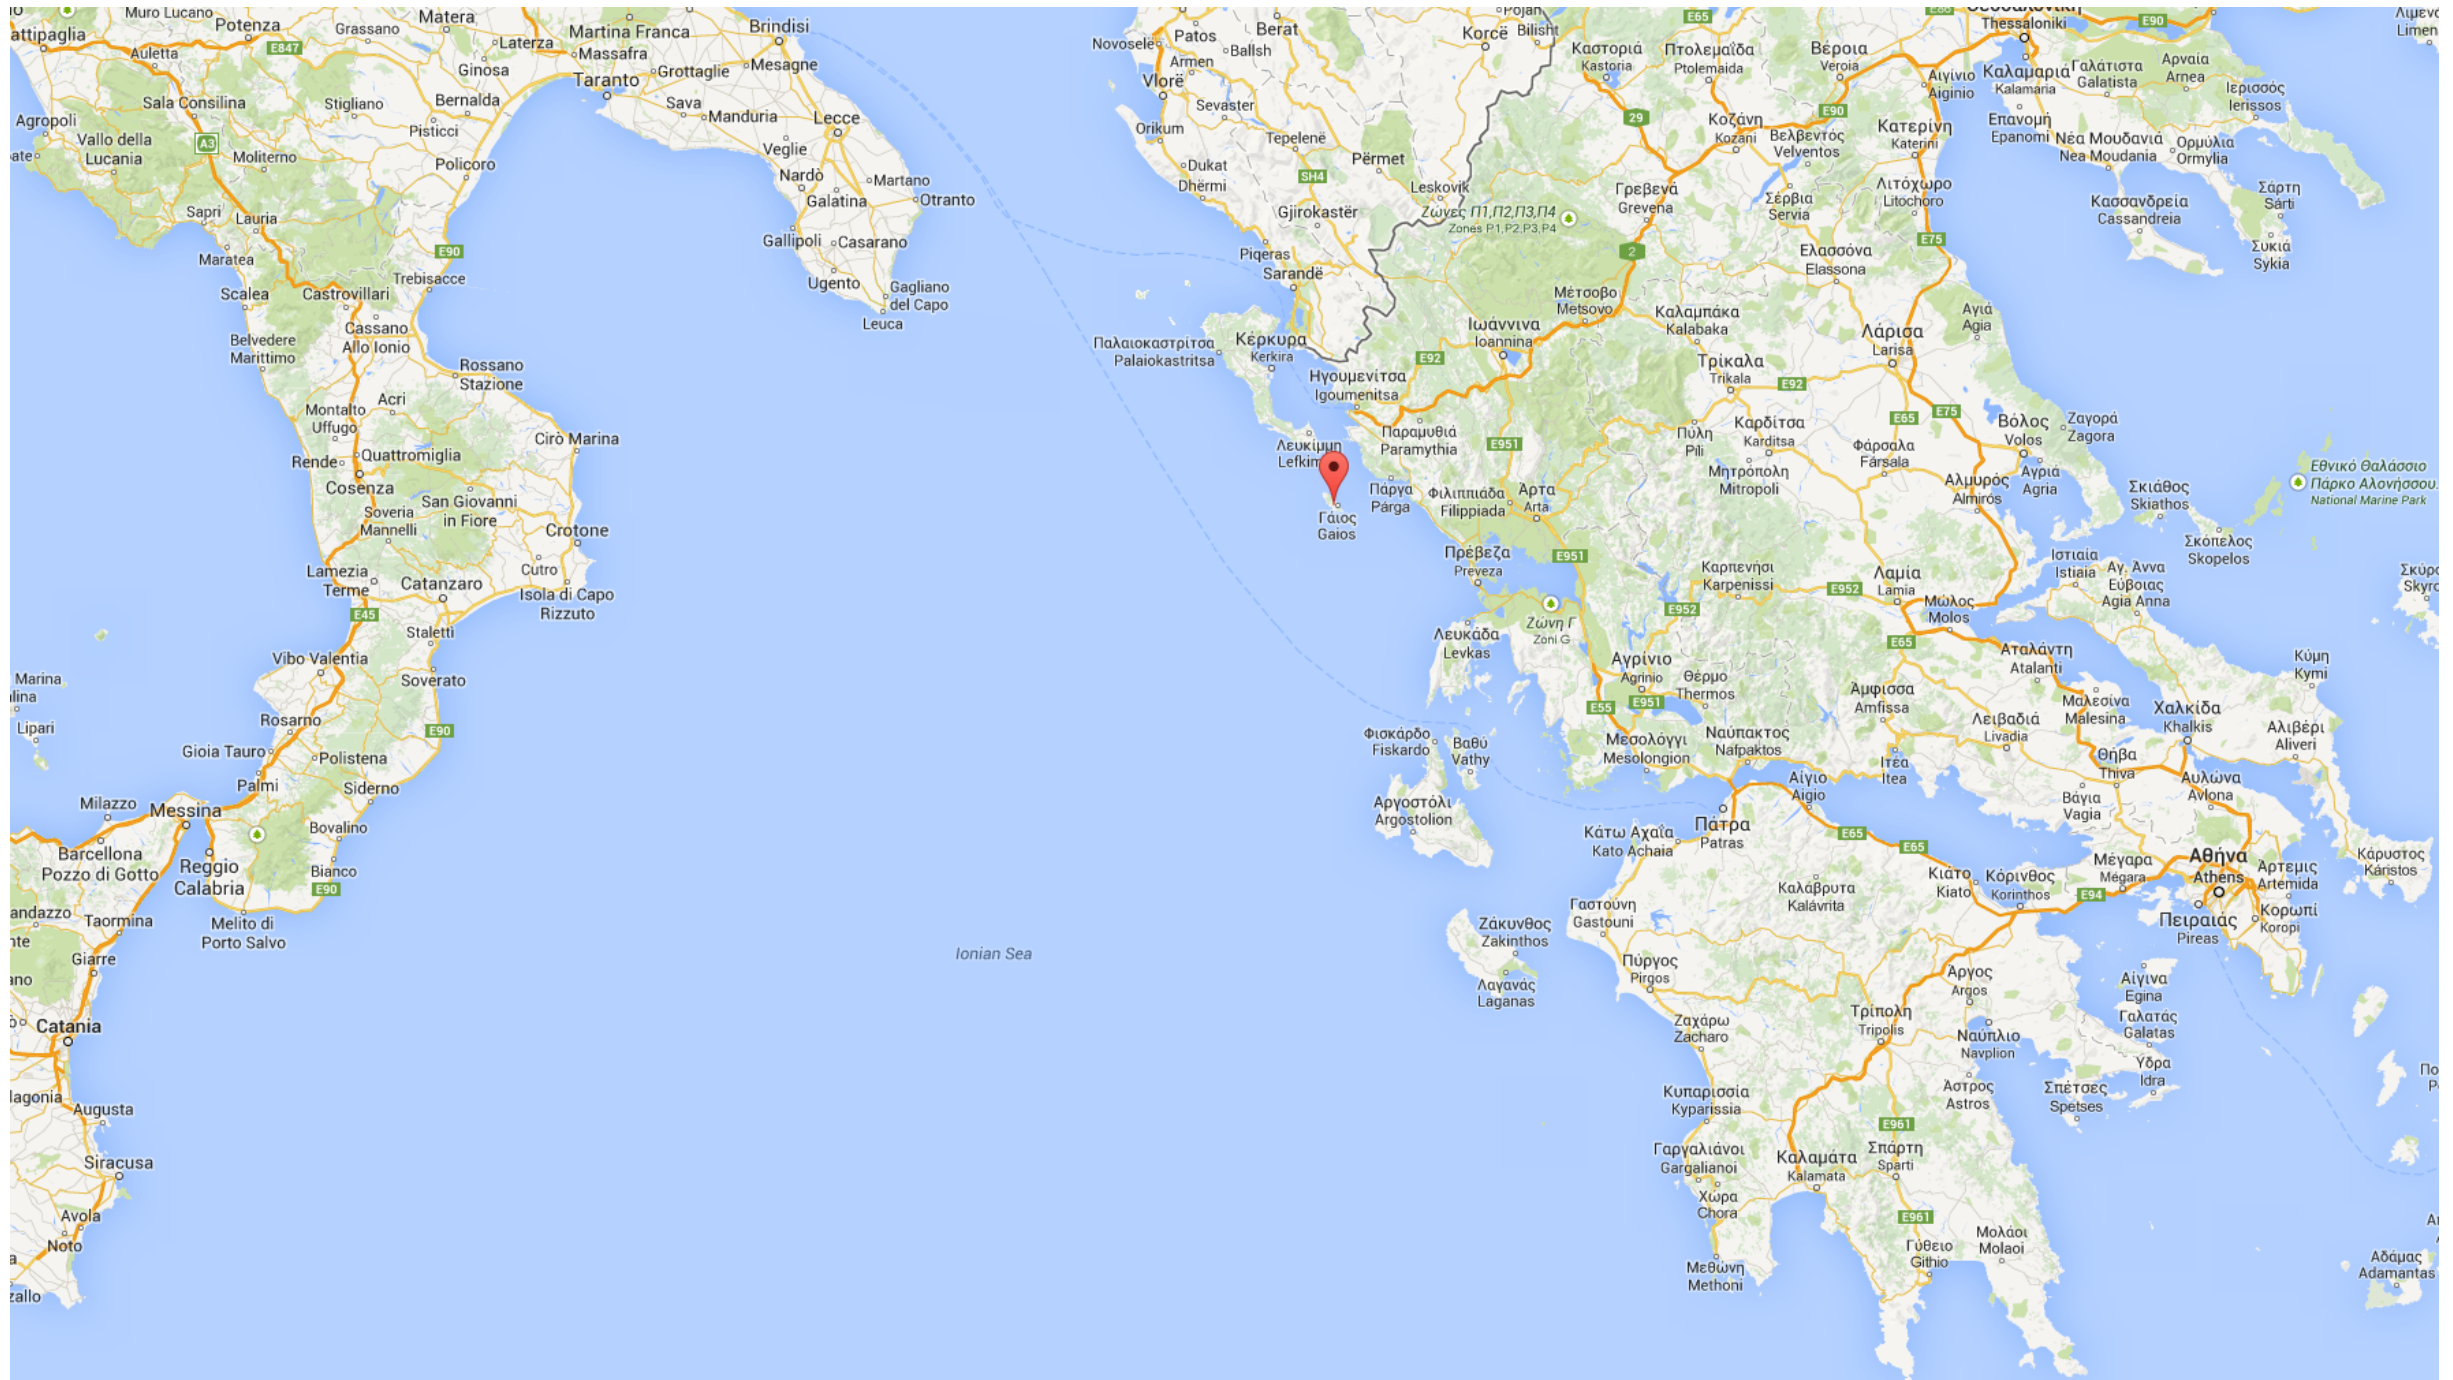


PAXOS



PAXOS

- Leslie Lamport. **The part-time parliament**. ACM Transactions on Computer Systems, 16(2):133–169, May 1998.
- Leslie Lamport described in 1990 the algorithm as the solution to a problem of the parliament on a fictitious Greek island called Paxos (not Italy)
- Many readers were so distracted by the description of the activities of the legislators, they did not understand the meaning and purpose of the algorithm. The paper was rejected.
- Leslie Lamport refused to rewrite the paper. He later wrote that he “was quite annoyed at how humorless everyone working in the field seemed to be”
- After a few years, some people started to understand the importance of the algorithm
- After eight years, Leslie Lamport submitted the paper again, basically unaltered. It got accepted!



Leslie Lamport
ACM Turing Award 2014

Consensus (again)

- We have a collection of processes
- Each process can propose a value
- A single value among the proposed values is chosen
- If no value is proposed, no value should be chosen
- If a value has been chosen, processes should be able to learn the chosen value

System Model

- Asynchronous communications
- Non-byzantine failures:
 - Nodes crash
 - Messages can be lost, duplicated arbitrarily late
 - No corrupted messages
- Fail-recover: crashed nodes may recover later

Roles


- Each nodes has one or more of 3 roles
- Proposer: a node that can propose a certain value for acceptance
- Acceptor: a node that receives proposals from proposers and that can either accept or reject a proposal
- Learner: a node not involved in the decision process that wants to know the final result of the decision process
- We will assume all nodes act both as proposer and as acceptor

Acceptors

- A single acceptor (a.k.a. coordinator) can fail and block the whole procedure
- There is not a coordinator, but multiple acceptors.
- An acceptor may **accept** a single value (e.g., express a single vote)
- How many acceptors do we need?

Majority

- To ensure that a single value is chosen, any majority (50% + 1) of acceptors is enough
- The intersection of two majorities is not empty
- An acceptor may **accept** a single value
- A value is **chosen** when a majority of acceptors has accepted it
- If a majority chooses a value, no other majority can choose a different value
- If an acceptor crashes, chosen value still available




Acceptors

Since there can be a single proposer, the algorithm must guarantee that

P1: An acceptor must accept the first proposed value it receives

Since there can be a single proposer, the algorithm must guarantee that a single value is chosen only when it is accepted by a majority of acceptors

Split Votes

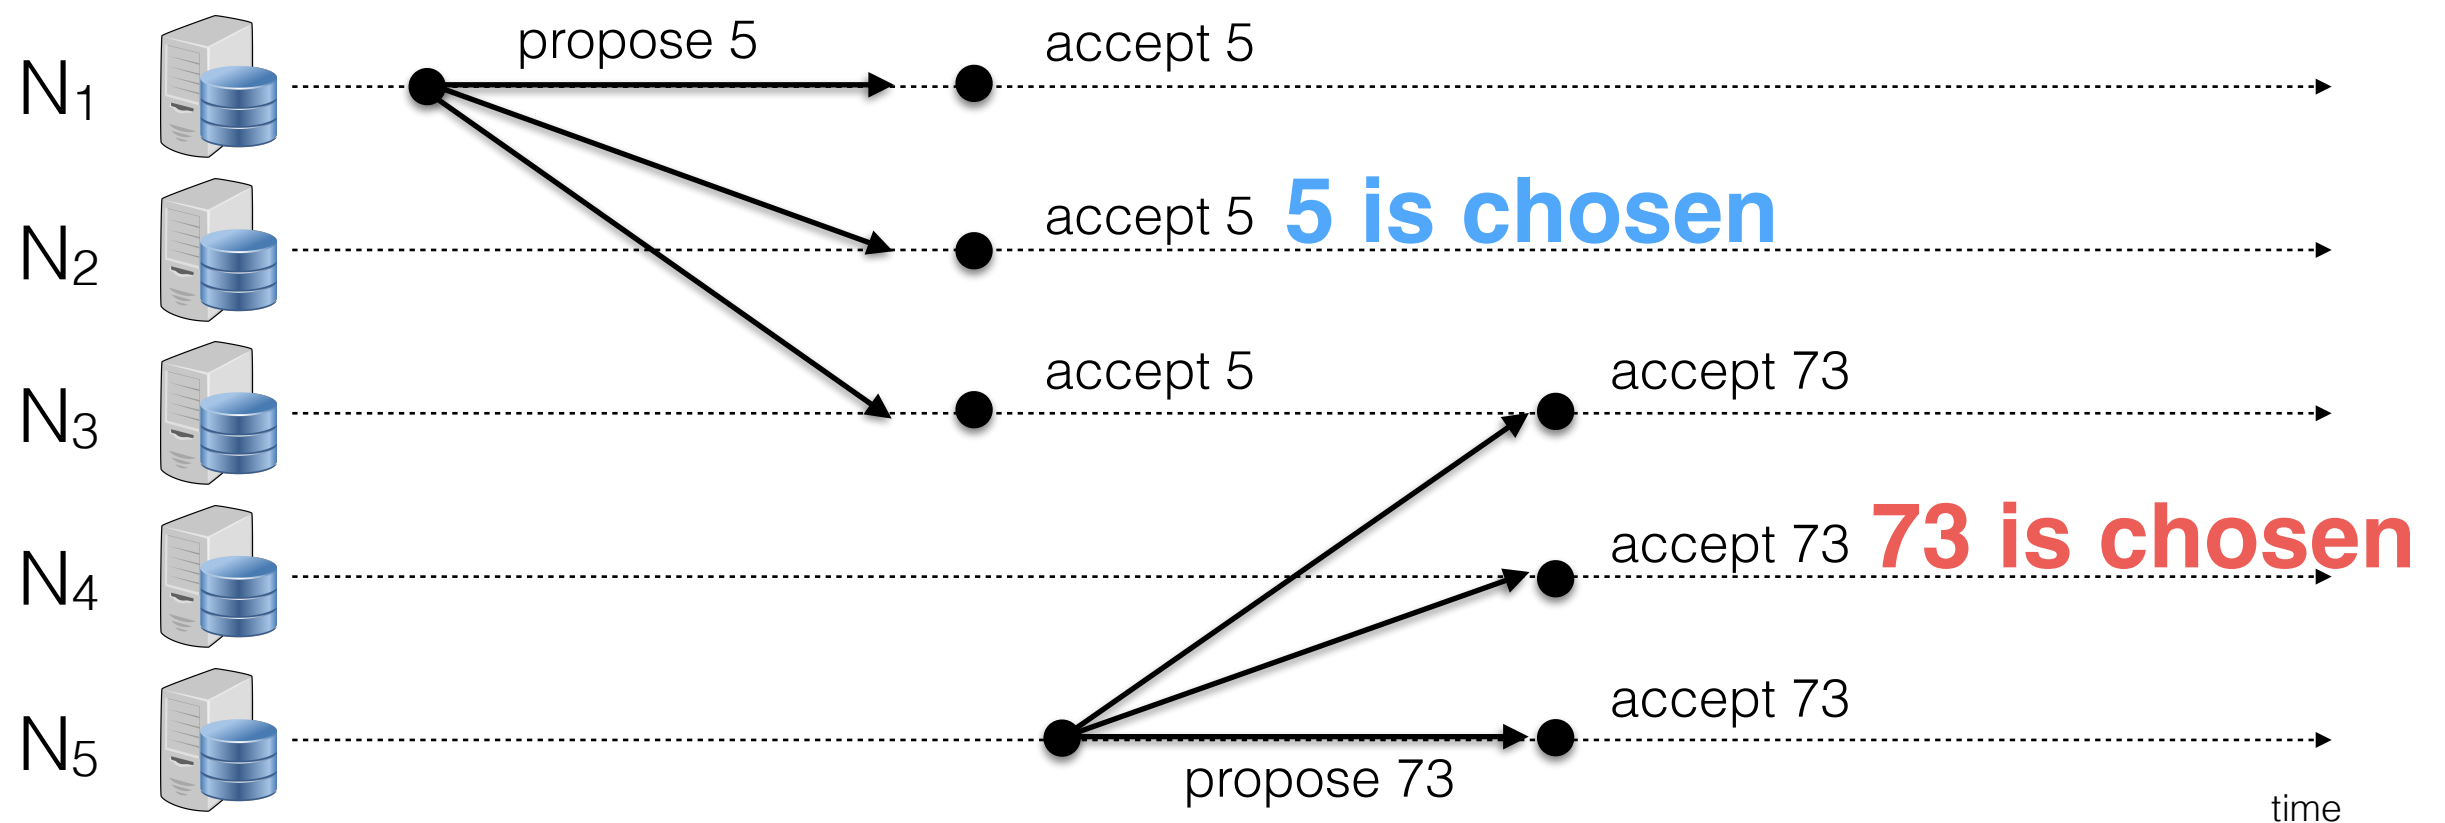


**An acceptor must accept the first proposed value it receives
(and can accept more than one proposed value)**

An acceptor can **accept** any number of proposed values, but an accepted proposed value may not necessarily be **chosen**.

Conflicting choices

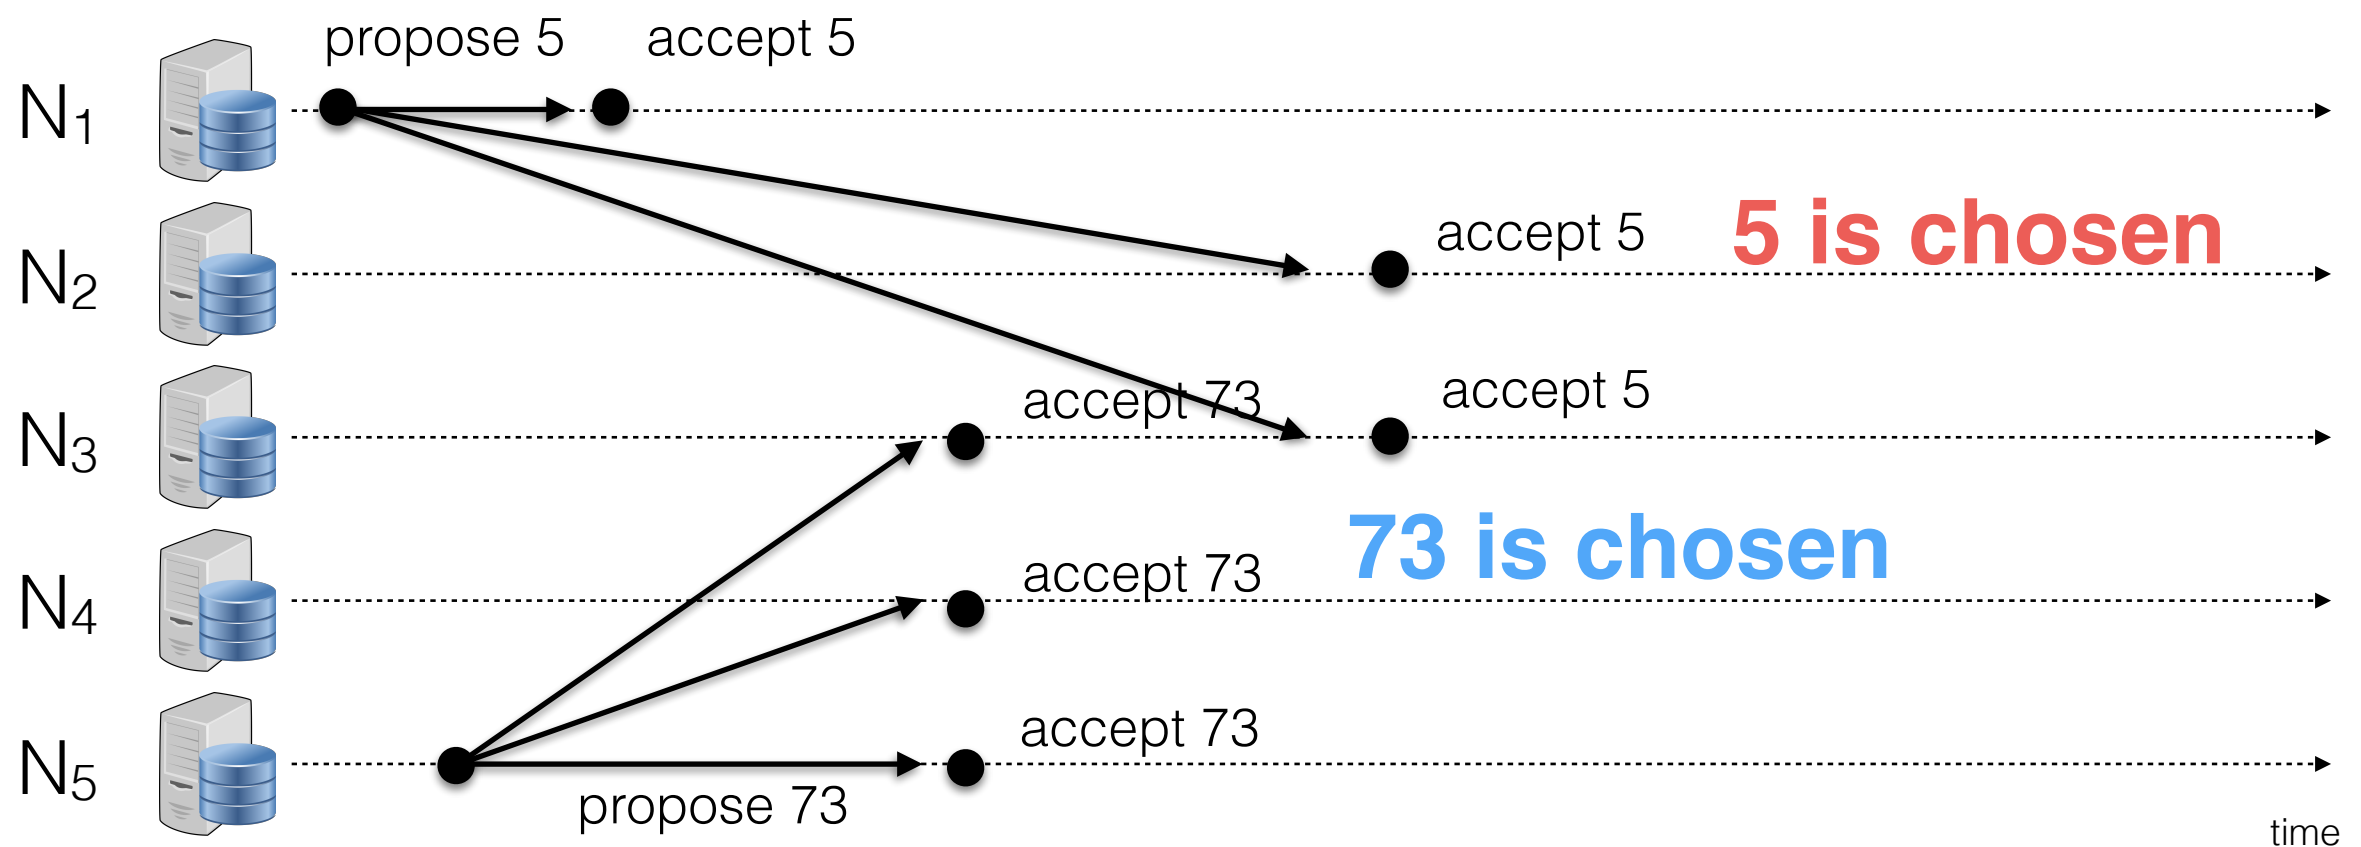
Should acceptors accept every proposed value? No



Once a value has been **chosen**, proposers must **propose** that same value

We need a two-phase protocol

Conflicting choices



Once a value has been **chosen**, older proposed values must be ignored

We need to order proposed values and reject old ones

Proposals

- A **proposal** (v, n) consists in the **proposed value** v and a **proposal number** n
- Whenever a proposer issues a new proposal, it chooses a strictly-increasing proposal number
- For a given proposer, a new proposal number must be greater than anything it has seen/used before
- Simple implementation



- Each server stores MAXROUND, i.e., the largest round number it has seen so far
- To generate a new proposal number: increment MAXROUND and concatenate with NODEID
- Proposers must persist MAXROUND on disk: must not reuse proposal numbers after crash/restart

Phase 1: PREPARE

- A proposer broadcasts a *prepare proposal* (v, n) with its own proposal value v and proposal number n
- An acceptor receives a *prepare proposal* (v, n):
 - if it has never received a *prepare proposal*:
 - it *promises* to never accept *proposal numbers* lesser than n ($n_{min} = n$)
 - it returns $(\emptyset, 0)$;
 - otherwise it checks its *promise*:
 - if $n > n_{min}$ it sends its last *accepted proposal* (v_{last}, n_{last}); note that $n_{last} < n$.
 - otherwise it does nothing

Phase 2: PROPOSE

- When a proposer receives a majority of responses, it broadcast to them a *propose proposal* (v' , n)
 - If all acceptors returned $(\emptyset, 0)$, v' is its own proposal value v
 - Otherwise v' is the v_{last} proposal value in the returned proposals with the greatest proposal number n_{last}
- If a majority of acceptors replies with **ACK**, the proposal is chosen
- Note that after a timeout, the proposer gives up and may send a new proposal

PAXOS Algorithm


Proposers

Acceptors

- 0) $n_p = 0$ highest prepare number seen
 $n_a = 0$ highest accepted proposal number
 $v_a = \emptyset$ highest accepted proposal value

These value must stably persist on disk

- 1) Choose new proposal number $n > n_p$
- 2) Broadcast PREPARE(v, n) to all nodes
- 4) If REPLY(v_a, n_a) from majority:
 $v' = v_a$ with greatest n_a
 If $v' = \emptyset$ then
 $v' = v$
- 5) Broadcast PROPOSE(v', n) to all nodes
- 7) If ACCEPT(v', n) from majority:
 Broadcast DECIDED(v') to all nodes



- 3) Handle PREPARE(v, n):
 If ($n > n_p$) then
 $n_p = n$
 REPLY(v_a, n_a)



- 6) Handle PROPOSE(v', n):
 If ($n \geq n_p$) then
 $n_p = n$
 $(v_a, n_a) = (v', n)$
 ACCEPT(v', n)



Learning a decision

- After a proposal is chosen, only the proposer knows about it!
- How do the other nodes get informed?
 1. The proposer could inform all nodes directly
 - If the proposer fails, the others are not informed (directly)...
 2. The acceptors could broadcast every time they accept a proposal
 - Much more fault-tolerant
 - Many accepted proposals may not be chosen...
 - Moreover, choosing a value costs $O(n^2)$ messages without failures!
 3. The proposer could inform some nodes directly
 - They will broadcast the decision to other nodes

PAXOS is safe!

If a proposal (v, n) is **chosen**, then for every proposed proposal (u, m) for which $m > n$ it holds that $v = u$

- Assume that there is a proposed proposal (u, m) for which $m > n$ and $u \neq v$. Consider such a proposal with the smallest m .
- Consider the non-empty intersection S of the two majority sets of nodes that are acceptors for (v, n) and (u, m) proposals.
- Since proposal (v, n) has been accepted and $m > n$, nodes in S must have received $\text{PREPARE}(u, m)$ after (v, n) has been accepted, thus returning $\text{REPLY}(v, n')$, with $n \leq n' < m$.
- As a consequence, the proposer of (u, m) should propose (v, m) , hence $u = v$, that is a contradiction.

PAXOS is correct!

If a value is chosen, all acceptors choose this value

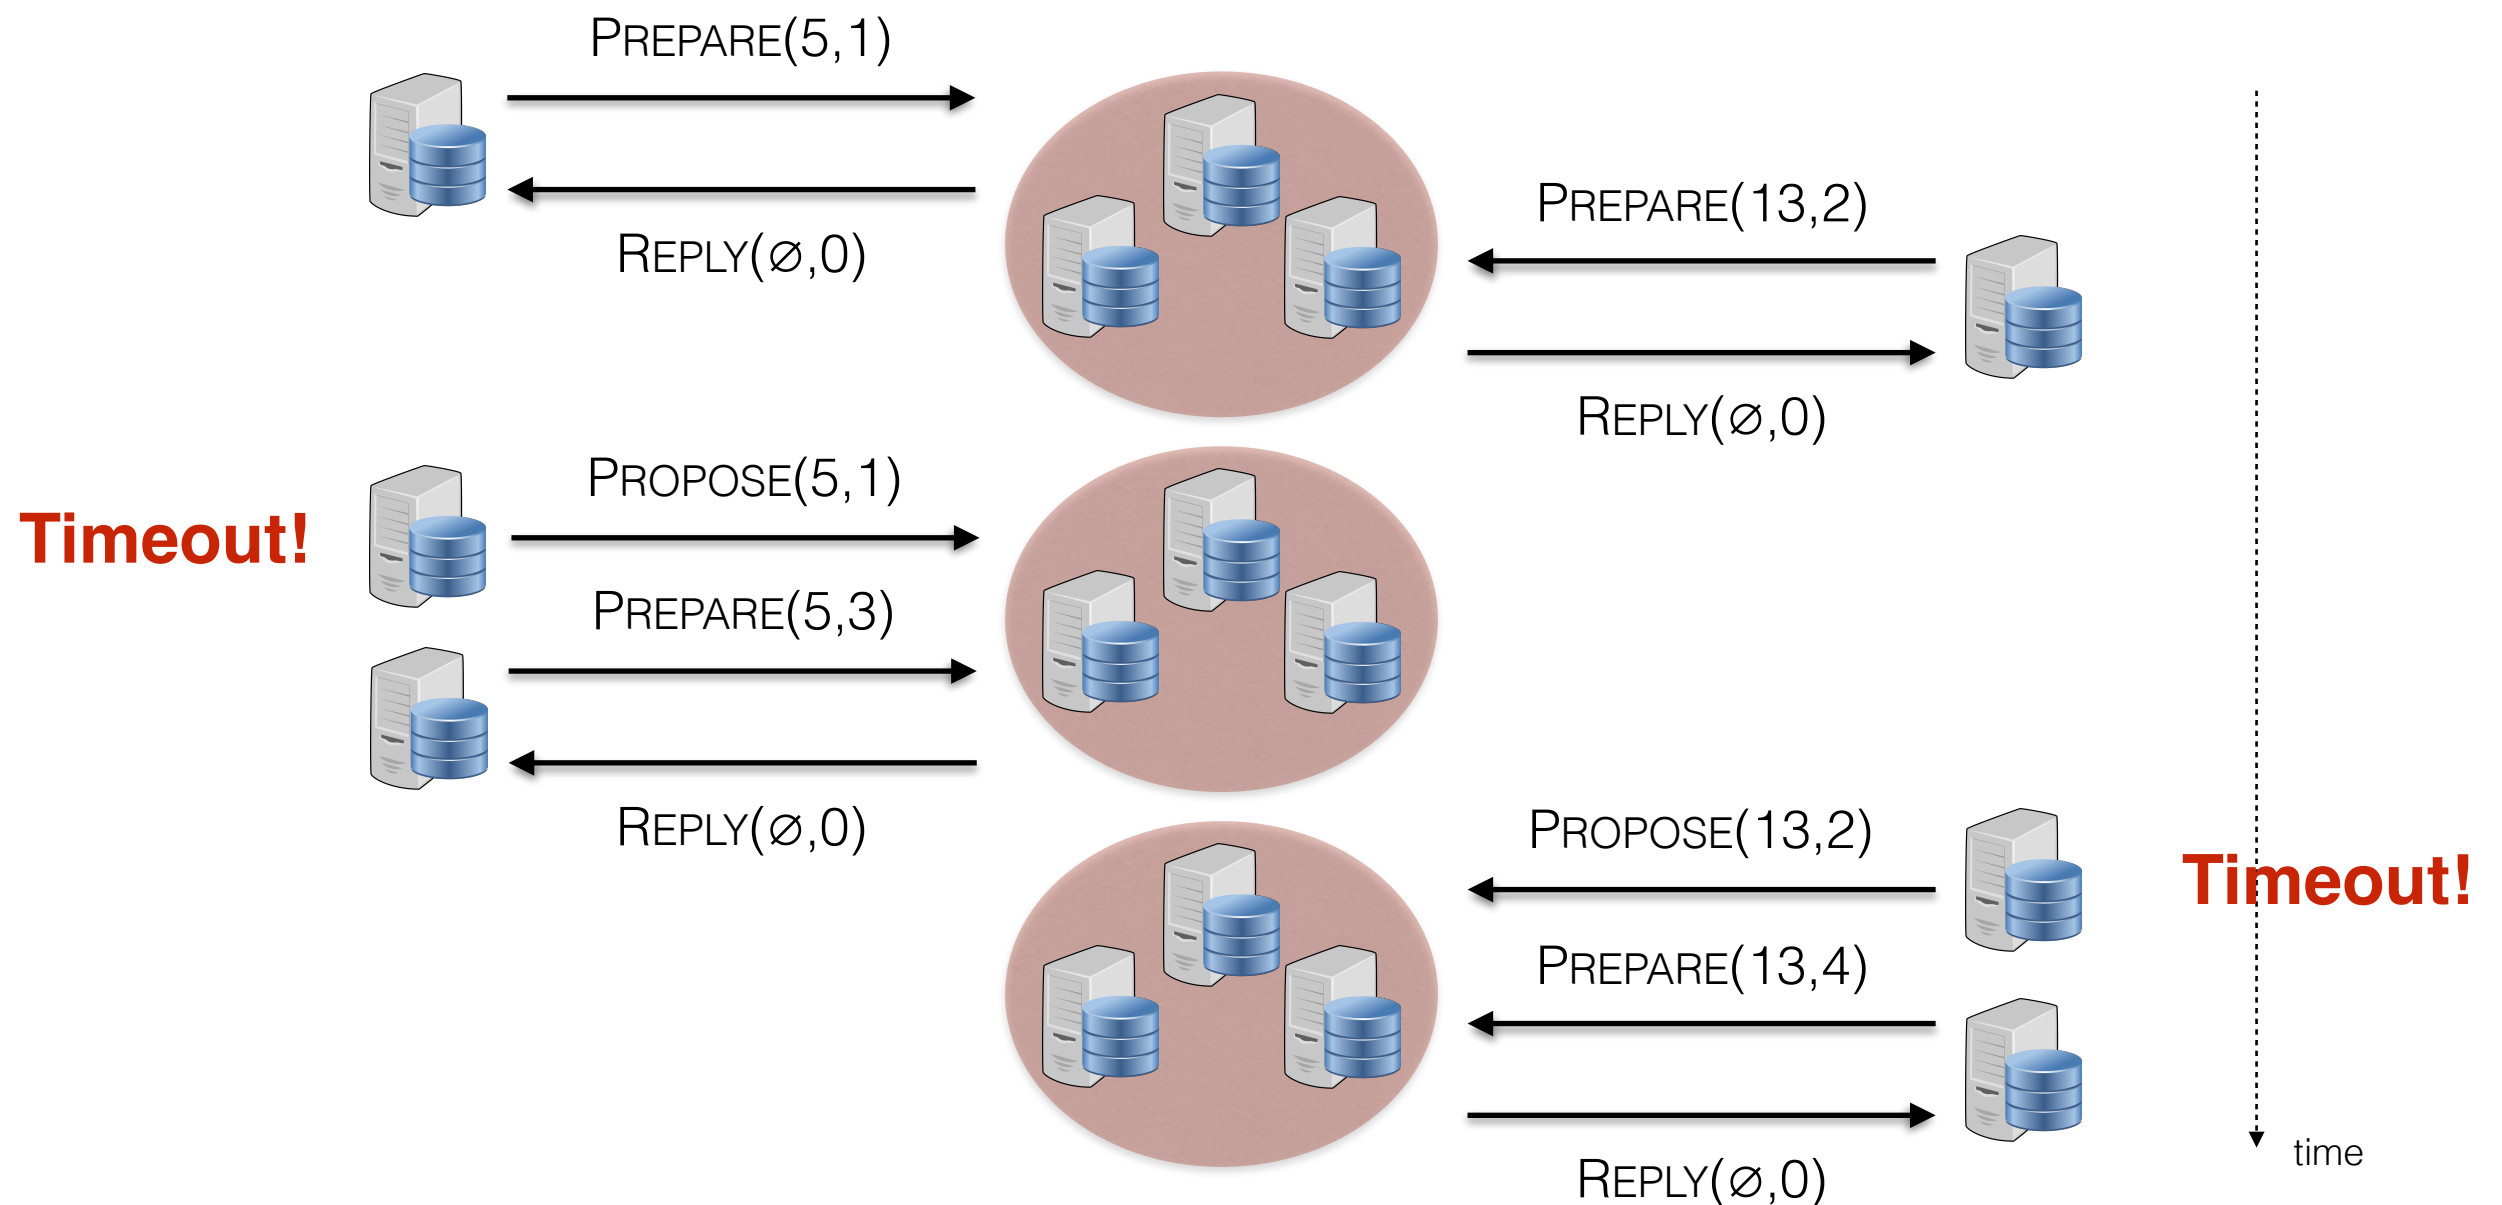
- Once a proposal (v, n) is chosen, each following proposal (u, m) has the same proposal value, i.e., $u = v$, according to the previous theorem.
- Since every following proposal has the same value v , every proposal that is accepted after (v, n) is chosen will have the same proposal value v .
- Since no other value than v can be accepted, no other values can be chosen.

PAXOS is great?

- PAXOS is a deterministic algorithm working for asynchronous systems and tolerating **$f < n/2$** failures.
- Many optimizations exists (Multi-PAXOS, Disk-PAXOS, RAFT)
- **FLP Theorem** [1985]. No totally correct consensus algorithm exists (for the given system model).
- PAXOS disproves the FLP Theorem?

No Liveness Guarantees

- PAXOS only guarantees that if a value is chosen, the other nodes can only choose the same value
- PAXOS does not guarantee that a value is chosen!



Correctness vs. Termination

- In asynchronous systems, we cannot guarantee termination and correctness at the same time
- PAXOS is correct, so termination is not guaranteed
- PAXOS cannot guarantee that a consensus is reached in a finite number of steps
- In practice, PAXOS can be optimized to reduce probability of no termination
- For example, the acceptors could send NAK if they do not accept a prepare message or a proposal (this optimization increases the message complexity)
- PAXOS is used in Apache's Zookeeper and Google's Chubby