

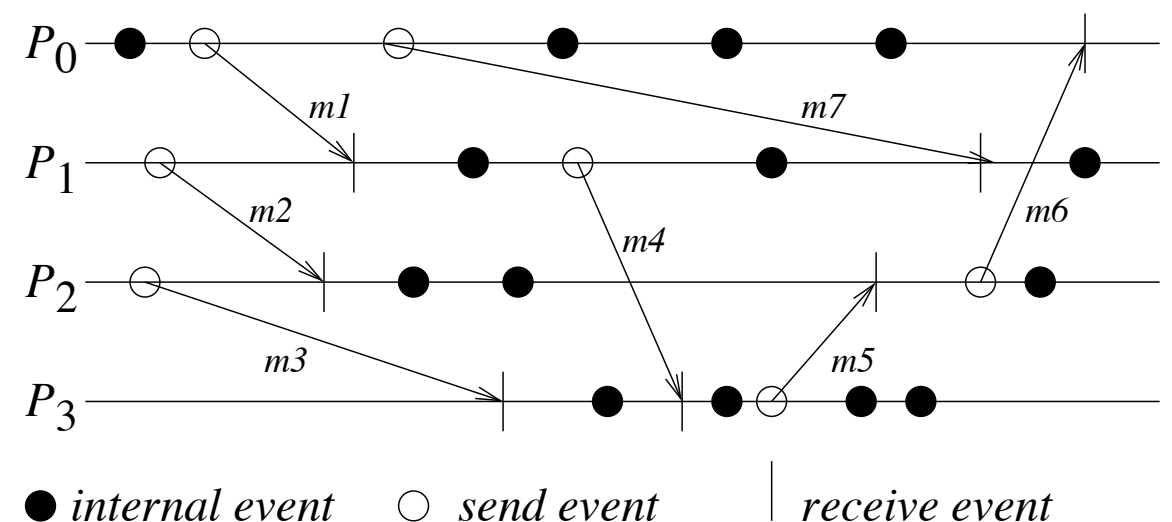
Logical Time

Causality and physical time

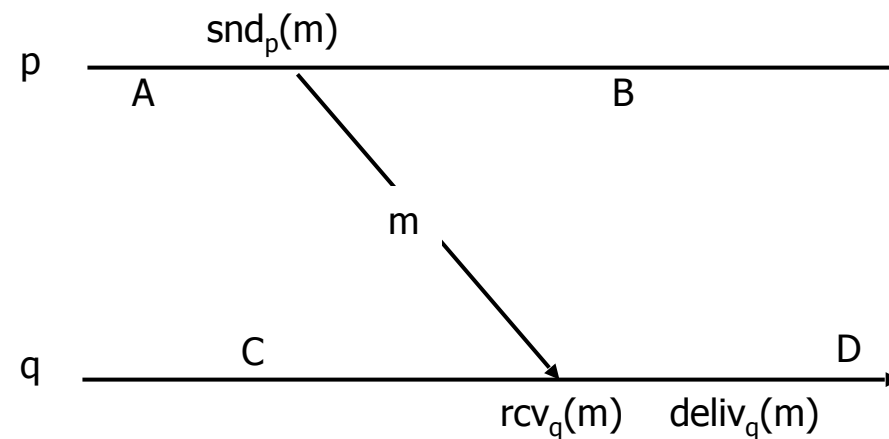
- **Causality** is fundamental to the design and analysis of parallel and distributed computing and OS.
 - Distributed algorithms design
 - Knowledge about the progress
 - Concurrency measure
- Usually causality is tracked using **physical time**.
- In distributed systems, it is **not possible** to have a **global physical time**, only an **approximation**.
 - **Network Time Protocol** (NTP) can maintain time accurate to a few tens of millisecond on the Internet
 - Not adequate to capture the causality relationship in distributed systems

Idea

- We **cannot sync** multiple clocks **perfectly**.
 - Thus, if we want to **order events** happened at different processes, we cannot rely on physical clocks.
- Then came **logical time**.
 - First proposed by Leslie **Lamport** in the 70's
 - Based on **causality** of events
 - Defined **relative time**, not absolute time
- **Critical observation**: time (ordering) only matters if two or more processes interact, i.e., send/receive messages.

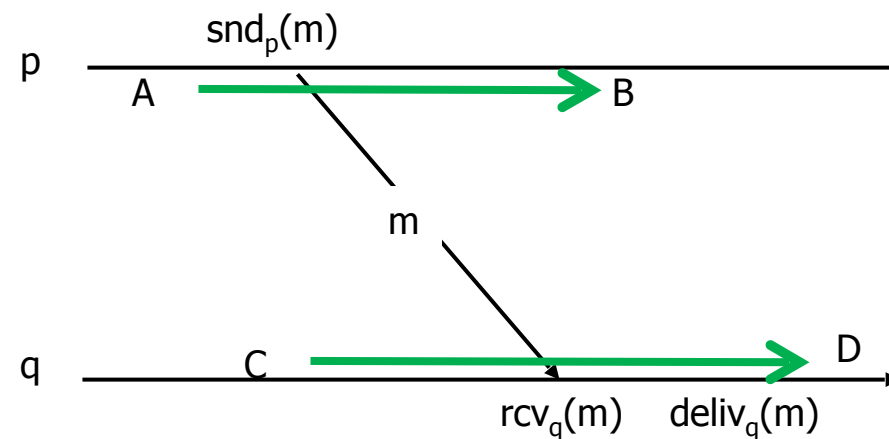


Time-line Diagrams



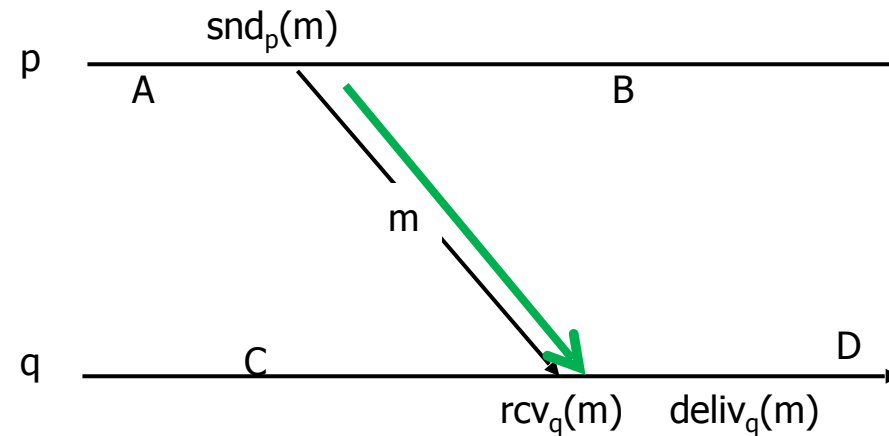
- A , B , C and D are “events”.
- Could be anything meaningful to the application
- So are $\text{snd}(m)$ and $\text{rcv}(m)$ and $\text{deliv}(m)$
- What ordering claims are meaningful?

Time-line Diagrams



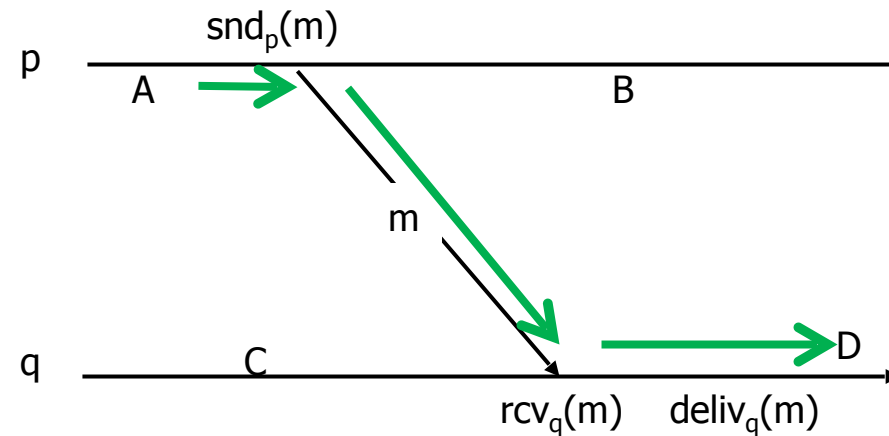
- A happens before B , and C before D
- “Local ordering” at a single process
- Write $A \rightarrow B$ and $C \rightarrow D$

Time-line Diagrams



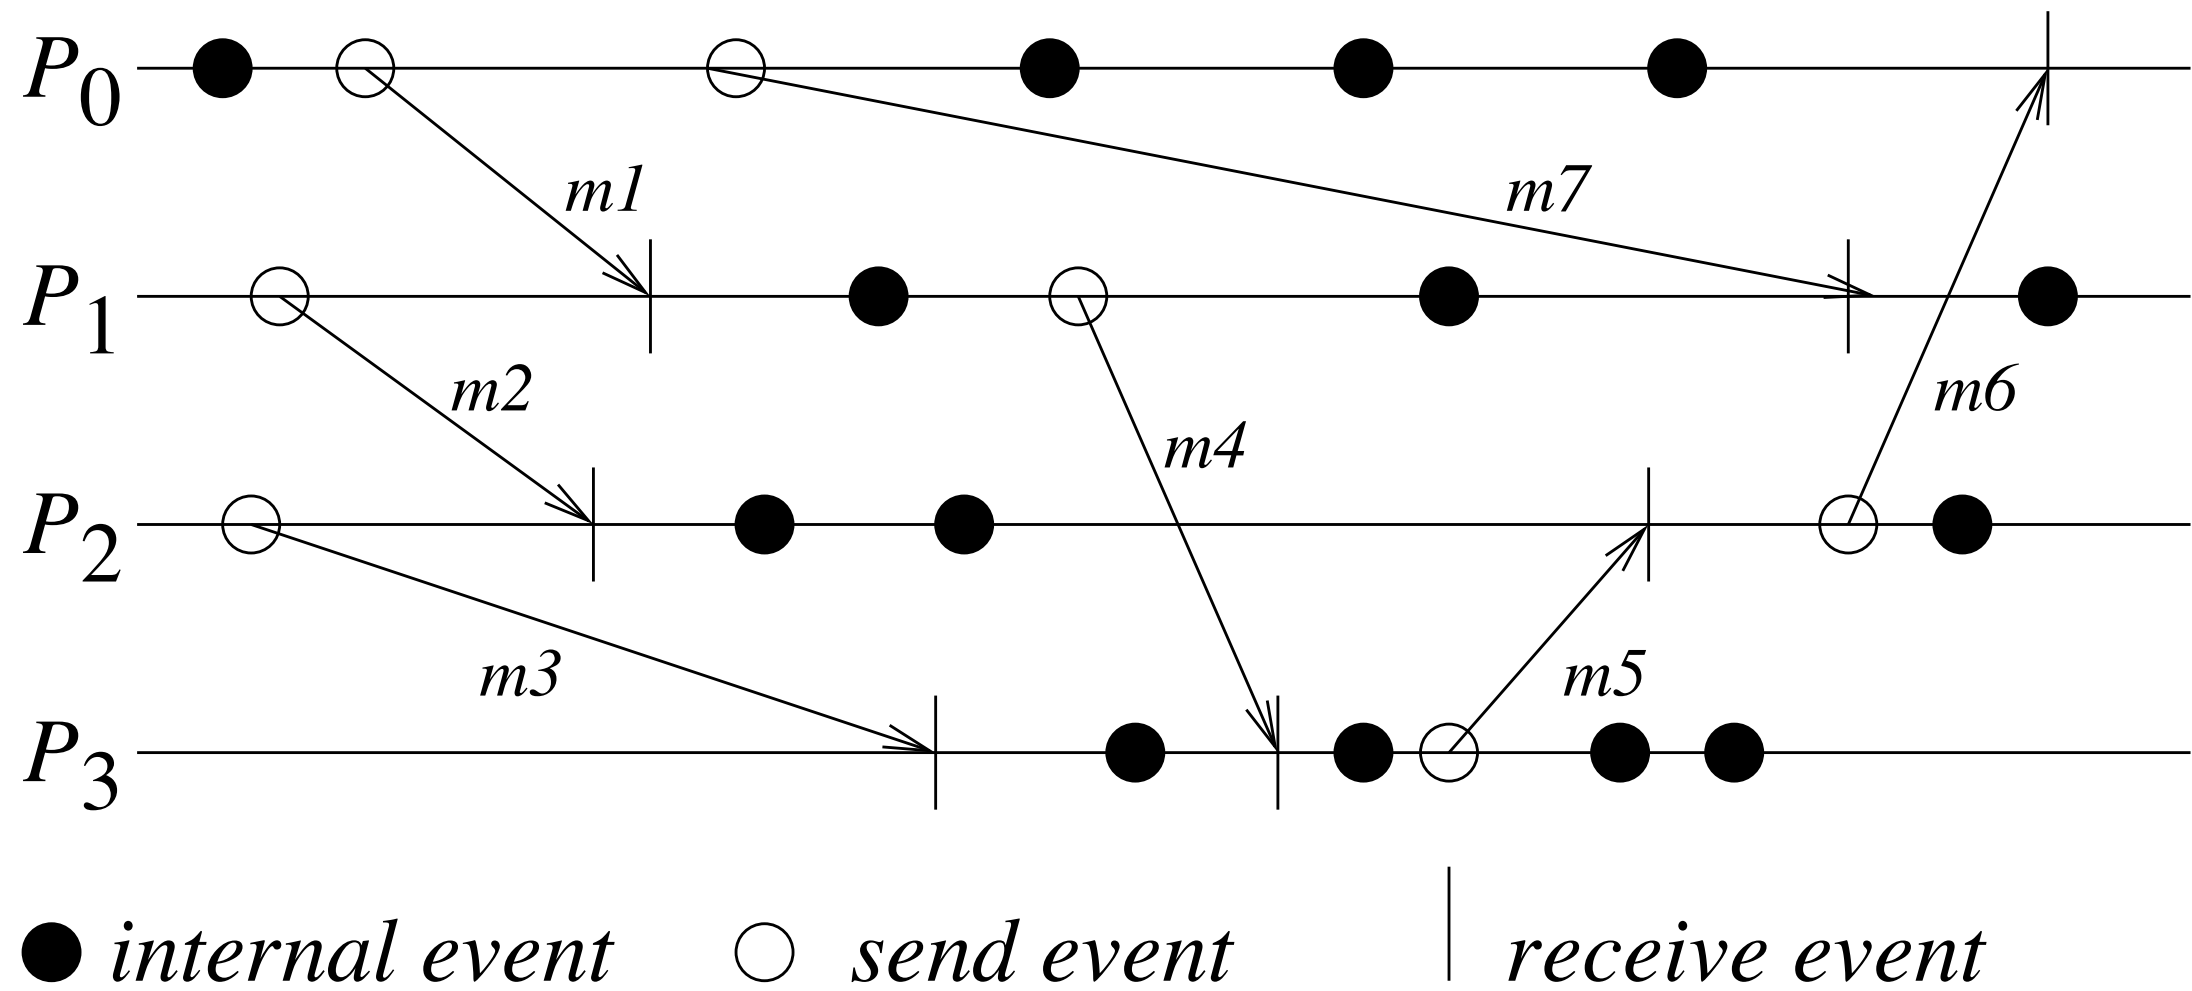
- $\text{snd}_p(m)$ also happens before $\text{rcv}_q(m)$
- “Distributed ordering” introduced by a message
- Write $\text{snd}_p(m) \rightarrow \text{rcv}_q(m)$

Time-line Diagrams



- A happens before D
 - Transitivity: A happens before $\text{snd}_p(m)$, which happens before $\text{rcv}_q(m)$, which happens before D
- B and D are concurrent
 - Looks like B happens first, but D has no way to know. No information flowed...

Events



Happens-Before Relation

- The **execution** of a distributed application results in a **set of** distributed **events** produced by the processes.
- Let **H** denote the set of events executed in a distributed computation.
- Define a **binary relation** on the set H, denoted as \rightarrow , that expresses **causal dependencies** between events in the distributed execution.
- \rightarrow is called **Happens-Before relation**.
- Properties:
 - On the same process: $a \rightarrow b$ if $\text{realtime}(a) < \text{realtime}(b)$
 - If p_1 sends m to p_2 : $\text{send}(m) \rightarrow \text{receive}(m)$
 - Transitivity: if $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$

System of Logical Clocks

- Informally:
 - Every process has a **logical clock** that is advanced according to some rules.
 - Every event is **assigned** a logical timestamp.
 - The \rightarrow relation between two events can be **inferred** from their timestamps.
 - Timestamps obey a **monotonicity property**: if $a \rightarrow b$, then $\text{timestamp}(a) < \text{timestamp}(b)$.
- Formally, a **system of logical clocks** is composed by:
 - a **time domain** T , whose elements form a partially ordered set over a relation $<$.
 - a **logical clock** C , that is a function mapping an event e in H to an element in the time domain T , denoted as $C(e)$ and called **timestamp** of e .
 - a logical clock C must satisfy the **clock consistency condition**:

$$\text{for two events } e_i \text{ and } e_j, \quad e_i \rightarrow e_j \Rightarrow C(e_i) < C(e_j)$$

- The system of clocks (T, C) is said to be **strongly consistent** if the following condition is satisfied:

$$\text{for two events } e_i \text{ and } e_j, \quad e_i \rightarrow e_j \Leftrightarrow C(e_i) < C(e_j)$$

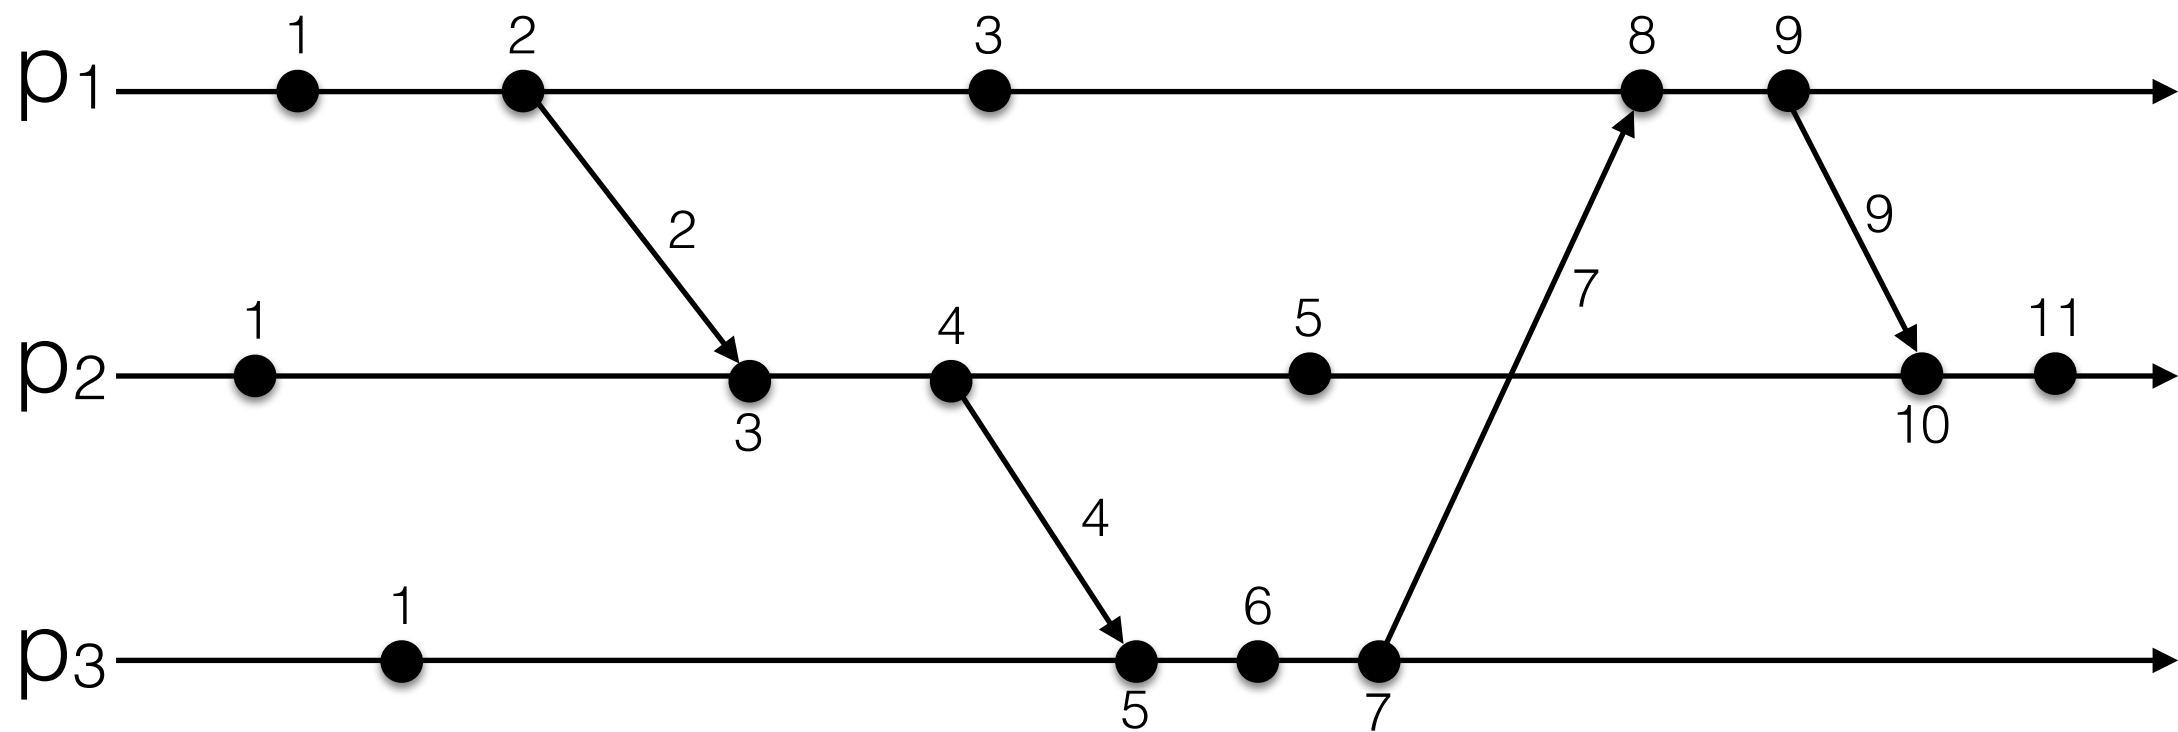
Implementation

- Implementation of logical clocks require:
 - **data structures** local to every process to represent logical time
 - a **set of rules** to update the data structures to ensure the consistency condition
- The **data structures** of a process p_i must allow it to:
 - measure its own progress, with a (**logical**) **local clock** lc_i
 - represent its own view of the logical global time to assign consistent timestamps to its local events, with a (**logical**) **global clock** gc_i
 - typically lc_i is a part of gc_i
- The rules must:
 - R1: decide how the logical local clock is updated by a process when it executes an event (send, receive, internal)
 - R2: decide how a process updates its logical global clock to update its view of the global time and global progress.


Scalar Clocks

- Proposed by Lamport in 1978.
- Time domain T is the set of **non-negative integers**.
- For each process p_i , the logical local clock and the logical global clock are squashed into **one integer variable** C_i .
- R1: before executing an event (send, receive, internal), process p_i executes the following:
$$C_i = C_i + d \ (d > 0)$$
 - In general every time R1 is executed, d can have a different value.
 - Typically d is kept at 1 to keep the rate of increase of C_i 's to its lowest values.
- R2: Each message piggybacks the clock value of its sender at sending time. When a process p_i receives a message with timestamp C_{msg} , it executes the following actions:
 1. $C_i = \max(C_i, C_{msg})$
 2. Execute R1
 3. Deliver the message to p_i

Example



Find the error...




Basic Properties

- The **consistency** property is **satisfied**.
- If $C(e_i) = C(e_j)$ then e_i and e_j are concurrent events.
- To **totally order** events, we need a **tie-breaking mechanism** for concurrent events. This is typically done by augmenting the scalar timestamp with a **process identifier**, e.g., (t,i) .
 - Process identifiers are linearly ordered and used to break ties.
- If $d=1$ we have that, if event e has a timestamp h , then $h-1$ represents the **minimum logical duration**, counted in units of events, required before producing event e .
- The **strong consistency** property is **NOT satisfied**.

Example

3 < 4 but the former did not happen before the latter



The lack of strong consistency is due to the squashing of logical local and global clocks into one

Vector Clocks (I)

- Proposed by **Fidge**, **Mattern** and **Schmuck** in 1988-1991.
- Time domain T is a set of n -dimension **non-negative integer vectors**.
- Each process p_i maintains a vector $vt_i[1..n]$.
- $vt_i[i]$ is the **logical local clock** of p_i .
- $vt_i[j]$ represents process p_i 's latest knowledge of process p_j local time. If $vt_i[j] = x$ then process p_i knows that local time at process p_j had progressed till x .

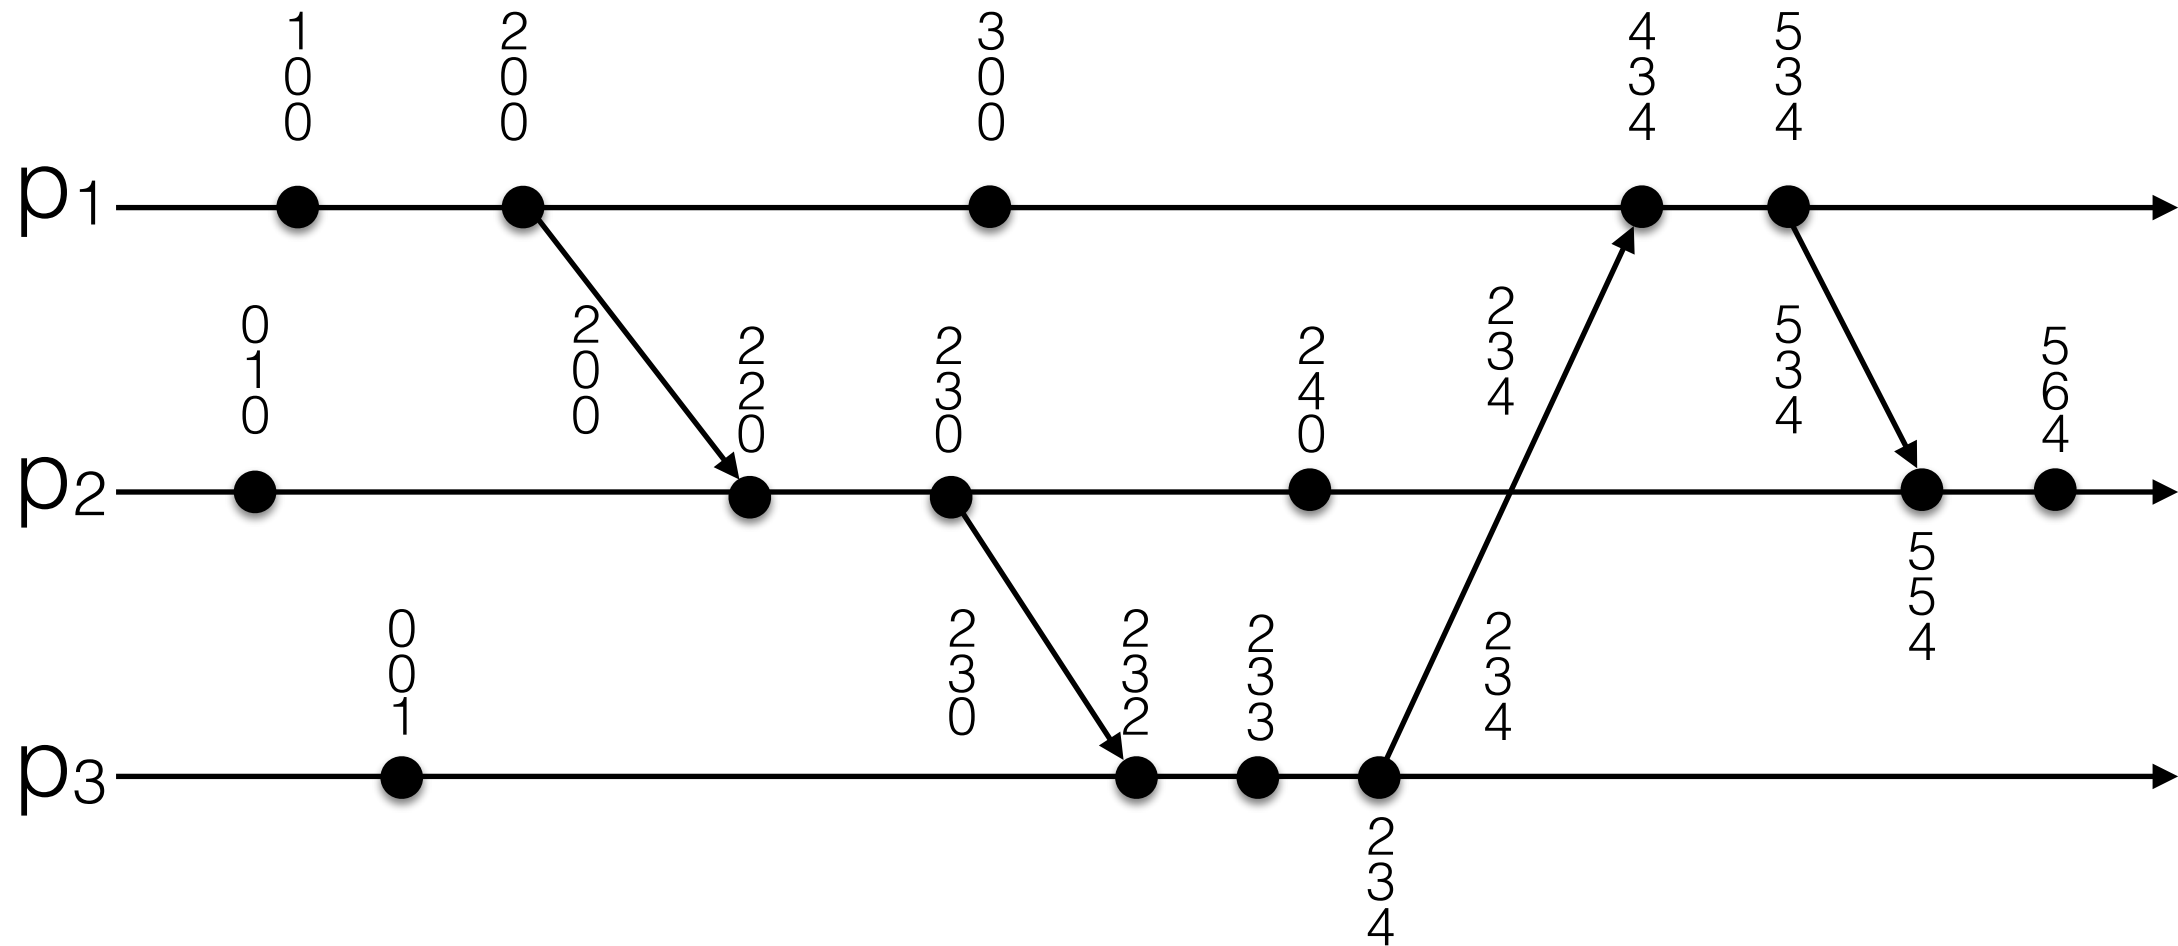
Vector Clocks (II)

- Initially $vt_i = [0, 0, 0, \dots, 0]$
- R1: before executing an event (send, receive, internal), process p_i executes the following:

$$vt_i[i] = vt_i[i] + d \ (d > 0)$$

- R2: Each message m is piggybacked with the vector clock vt of the sender process at sending time. When a process p_i receives a message with (m, vt) , it executes the following actions:
 1. Update its logical global time as follows:
$$1 \leq k < n: vt_i[k] = \max(vt_i[k], vt[k])$$
 2. Execute R1
 3. Deliver the message m to p_i

Example



Comparing Vector Clocks

- $VT_1 = VT_2$
 - iff $VT_1[i] = VT_2[i]$, for all $i = 1, \dots, n$
- $VT_1 \leq VT_2$,
 - iff $VT_1[i] \leq VT_2[i]$, for all $i = 1, \dots, n$
- $VT_1 < VT_2$,
 - iff $VT_1 \leq VT_2$ & $\exists j (1 \leq j \leq n \text{ \& } VT_1[j] < VT_2[j])$
- $VT_1 \parallel VT_2$
 - iff $\neg(VT_1 \leq VT_2) \text{ \& } \neg(VT_2 \leq VT_1)$

Basic Properties

- The **consistency** property is **satisfied**.
- The **strong consistency** property is **satisfied** (using always at least n elements).
- If two events x and y have timestamps v_h and v_k respectively, then we have the following **isomorphism**:

$$x \rightarrow y \Leftrightarrow v_h < v_k$$


$$x \parallel y \Leftrightarrow v_h \parallel v_k$$

- If $d = 1$ then we have the **event counting** property of scalar clocks for logical local clocks.
- Since vector clocks are strongly consistent they can **track causal dependencies exactly**.

Temporal Distortions


- Things can be complicated because we can't predict
 - Message delays (they vary constantly)
 - Execution speeds (often a process shares a machine with many other tasks)
 - Timing of external events

What does "now" mean?



- Timelines can “stretch”...
 - ... caused by scheduling effects, message delays, message loss...
- Timelines can “shrink”
 - E.g. something lets a machine speed up

Cuts




- Cuts represent instants of time.
 - But not every “cut” makes sense
 - Red cuts could occur but not gray ones.

Consistent Cuts

- Idea is to identify system states that “might” have occurred in real-life
- Need to avoid capturing states in which a message is received but nobody is shown as having sent it
- This is the problem with the gray cuts

Inconsistent Cuts



- Red messages cross grey cuts "backwards" in time

Consistent Cuts and Snapshots

- We want to draw a line across the system state such that
 - Every message “received” by a process is shown as having been sent by some other process
 - Some pending messages might still be in communication channels
- And we want to do this while running

System Assumptions

- Communication must be reliable
- Processes be failure-free.
- Point-to-point message delivery must be ordered (FIFO channels)
- Process/channel graph must be strongly connected
- No global shared memory and physical clock
- Under these assumptions, and after the algorithm completes, each process hold a copy of its local state and a set of messages that was in transit, with that process as their destination, during the snapshot.

System Assumptions

- At any instant, the state of process p_i , denoted by LS_i , is a result of the sequence of all the events executed by p_i till that instant.
- For an event e and a process state LS_i , $e \in LS_i$ iff e belongs to the sequence of events that have taken process p_i to state LS_i .
- For an event e and a process state LS_i , $e \notin LS_i$ iff e does not belong to the sequence of events that have taken process p_i to state LS_i .
- For a channel C_{ij} , the following set of messages can be defined based on the local states of the processes p_i and p_j

$$\text{transit}(LS_i, LS_j) = \{ m_{ij} \mid \text{send}(m_{ij}) \in LS_i \text{ and } \text{recv}(m_{ij}) \notin LS_j \}$$

Consistent Global State

- The **global state** of a distributed system is a collection of the local states of the processes (LS_i) and the channels (SC_{ij}).
- Notationally, global state GS is defined as

$$GS = \{ \bigcup_i LS_i, \bigcup_{i,j} SC_{ij} \}$$

- A global state GS is a **consistent** global state iff it satisfies the following two conditions :
 - **C1**: $\text{send}(m_{ij}) \in LS_i \Rightarrow m_{ij} \in SC_{ij} \text{ XOR } \text{recv}(m_{ij}) \in LS_j$
 - **C2**: $\text{send}(m_{ij}) \notin LS_i \Rightarrow m_{ij} \notin SC_{ij} \text{ AND } \text{recv}(m_{ij}) \notin LS_j$

Issues

- How to distinguish between the messages to be recorded in the snapshot from those not to be recorded?
 - Any message that is sent by a process before recording its snapshot, must be recorded in the global snapshot (from C1).
 - Any message that is sent by a process after recording its snapshot, must not be recorded in the global snapshot (from C2).
- How to determine the instant when a process takes its snapshot?
 - A process p_j must record its snapshot before processing a message m_{ij} that was sent by process p_i after recording its snapshot.

Chandy-Lamport Algorithm

- The algorithm uses a control message, called a **marker** whose role in a FIFO system is to separate messages in the channels.
- After a site has recorded its snapshot, it sends a **marker**, along all of its outgoing channels before sending out any more messages.
- A **marker** separates the messages in the channel into those to be included in the snapshot from those not to be recorded in the snapshot.
- A process must record its snapshot no later than when it receives a marker on any of its incoming channels.

Chandy-Lamport Algorithm

- The algorithm can be initiated by any process by executing the **Marker Sending Rule** by which it records its local state and sends a marker on each outgoing channel.
- A process executes the **Marker Receiving Rule** on receiving a marker. If the process has not yet recorded its local state, it records the state of the channel on which the marker is received as empty and executes the **Marker Sending Rule** to record its local state.
- The algorithm terminates after each process has received a marker on all of its incoming channels.
- All the local snapshots get disseminated to all other processes and all the processes can determine the global state.

Rules

- **Marker Sending Rule** for process p_i
 1. Process p_i *records* its state.
 2. For each outgoing channel C on which a marker has not been sent, p_i *sends* a marker along C before p_i sends further messages along C .
- **Marker Receiving Rule** for process p_j
 1. On *receiving* a marker along channel C :
if p_j has not *recorded* its state then
 1. *Record* the state of C as the *empty set*
 2. Follow the **Marker Sending Rule**
 - else
 1. *Record* the state of C as the *set of messages* received along C after p_j 's state was recorded and before p_j received the marker along C

Correctness & Complexity

- **Correctness**

- Due to FIFO property of channels, it follows that no message sent after the marker on that channel is recorded in the channel state. Thus, *condition C2 is satisfied.*
- When a process p_j receives a message m_{ij} that precedes the marker on channel C_{ij} , it acts as follows: if process p_j has not taken its snapshot yet, then it includes m_{ij} in its recorded snapshot. Otherwise, it records m_{ij} in the state of the channel C_{ij} . Thus, *condition C1 is satisfied.*

- **Complexity**

- The recording part of a single instance of the algorithm requires $O(e)$ messages and $O(d)$ time, where e is the number of edges in the network and d is the diameter of the network.